


CS3102 P2: Practical Report

Reliable Data Transfer Using UDP



University of
St Andrews

190010906

01 April 2022

1 Introduction

Reliable Data Transport (RDT) is a unicast, connection-oriented that provides reliable, ordered data transfer. This report will detail, analyse and evaluate the design and implementation of RDT.

2 Design

Given the scope of the practical, simplicity was the main goal when designing the RDT protocol. This section will discuss design decisions and the rationale behind them. The two attempted extension features, checksums and adaptive re-transmission timeouts, are also detailed.

2.1 Packet Structure

RDT packets (see Figure ??) are composed of a constant 12 byte header and an optional data segment. The data segment size ranges from 0 to 1300 bytes, with 1300 bytes used as the maximum size to allow testing with `slurpe-3` without issue. Given that the TCP Maximum Segment Size (MSS) [?] is generally 1500 bytes, this was considered a reasonable choice. Larger packet sizes may also lead to fragmentation at the IP Layer, which was undesirable.




Figure 1: RDT Packet Structure

The RDT header (see Figure ??) is comprised of the following fields: a 32-bit **sequence** field, used for keeping track of the number of bytes sent by the client; a 16-bit **type** field, used to denote the packet function; a 16-bit **checksum** field, calculated over the header and data segment to detect bit-errors; a 16-bit **size** field denoting the size of the data segment (in bytes); and a 16-bit **padding** field to ensure 32-bit word alignment.

Several factors influenced the RDT header design. As file sizes are often calculated as 32-bit integers (such as by the C library function `ftell`), a 32-bit **sequence** field was required to support the transmission of large files/amounts of data. The given implementation of the IPv4 header checksum returns a `uint16_t` value, thus necessitating a 16-bit **checksum** field. As a maximum data segment size of 1300 bytes was required, at least 11 bits were required for the **size** field; however 16 bits were used for alignment. There were no other considerations for field size other than 32-bit alignment for the **type** and **padding** fields.

For simplicity, a single **type** field was chosen rather than a set TCP-style flags. Given the minimal nature of the RDT protocol, it was faster simpler to enumerate all packet types (see below), rather than testing multiple flags.




Figure 2: RDT Header

The type field supports the following types: **SYN** (0) and **SYN ACK** (1), used for the connection handshake; **DATA** (2) and **ACK** (3), used for sending and acknowledging data segments; **FIN** (4) and **FIN ACK** (5), used for graceful connection termination; and **RST** (6), used for abrupt connection termination.

2.2 Connection Management

The operation of the RDT protocol is modelled by the FSM in Figure ???. For connection management, a two-way handshake is used. As RDT only supports uni-directional communication, a two-way handshake (see Figure ??) is adequate for establishing and terminating connections. Adaptive re-transmission timeouts are used in both the handshakes and transmission of data segments.

During handshakes, timeouts are used, with an initial RTO value of 200ms. This is doubled successively (Section ??). After five timeouts, the client with aborts with an **RST** to prevent clients from waiting on unavailable hosts or a dropped **FIN ACK**.

2.3 Data Transfer

RDT uses an Idle-RQ mechanism for sending data segments and will wait for acknowledgement of each sent segment. If the sender receives an **ACK** with a sequence number that is greater than expected or an RTO is triggered, it will retransmit the same packet. If it receives an **ACK** with a sequence number lower than expected, it will retransmit from that segment as it assumes a data segment has been dropped or corrupted...

For the handshake, a random sequence number is chosen to prevent ‘old’ or ‘stale’ packets from causing errors and to reduce predictability that could be exploited. The method for generating initial sequence numbers is pseudo-random and not cryptographically secure.

2.4 Adaptive RTO

As an extension, adaptive re-transmission timeouts using measured RTT has been implemented. RDT’s adaptive RDT is modelled on TCP’s Re-Transmission Timer [?]¹, with the same initial RTO of 1s and a maximum of 60s.

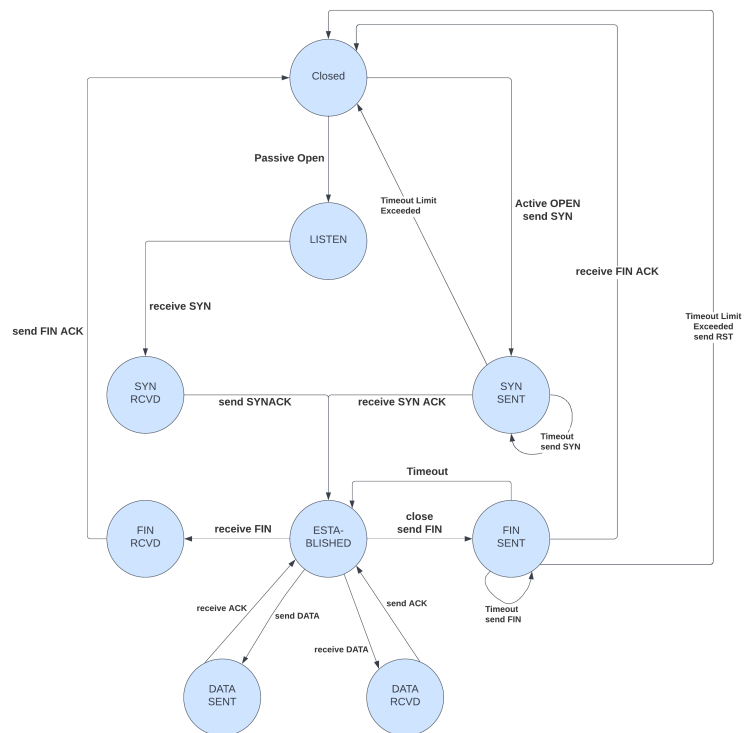


Figure 3: RDT Finite State Machine (see also A.1)

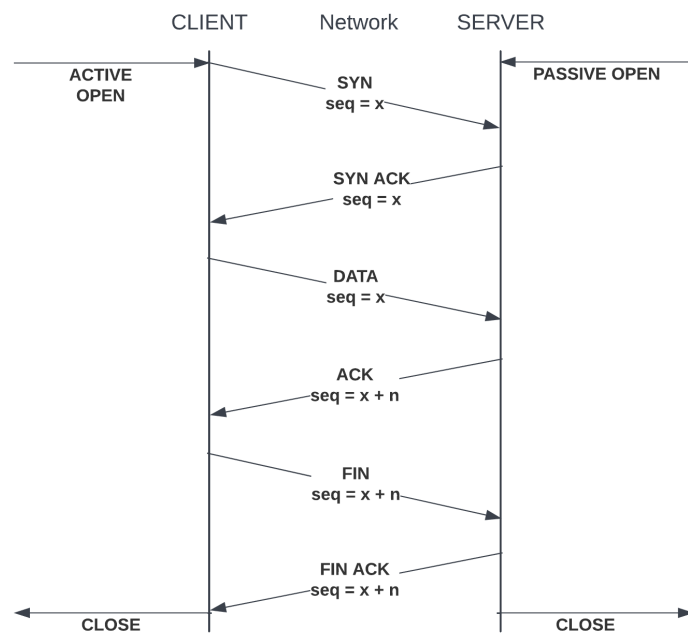


Figure 4: Connection Timeline Diagram (see also A.2)

2.5 Checksums

The RDT `checksum` field is calculated over the entire segment/packet (i.e. header and data) using the IPv4 Header Checksum Algorithm [?]. The implementation of the original source [?] to support the use of 8-bit byte values in the data segment. During checksum calculating, the `checksum` field itself is set to 0 for consistency.

3 Testing

To test and verify the correct operation of RDT, two test programs `RdtClient.c` and `RdtServer.c`, were created to send an arbitrary file between `pc7-082-1` and `pc7-085-1` using RDT. For test cases where delay, loss and a constrained data rate was introduced, `slurpe-3` was run on `pc7-087-1`.

3.1 Methodology

To test that RDT packets were delivered reliably and in-order, a 230KB JPEG file (*dog.jpg*²) was used as a test file. This provided visible feedback as to the integrity of the data received by `RdtServer.c`. This integrity was also verified by calculating SHA1 checksums of both the sent and received files. `slurpe-3` was used to provide an emulated path with loss, delay and restricted data rate to test the reliability of RDT in degraded network conditions.

To collect RTT data for analysing the performance of RDT, two test programs, `RdtClientRTT.c` and `RdtServerRTT.c`, were created. These programs performed 10 rounds of transferring the same 230 KB file to calculate a sample of 10 average RTT values. This was performed without `slurpe-3` to give a true characterisation of performance. These results are included in `data/rtt-results.csv`.

3.2 Results

From the results in Table ?? and Appendix ??, we can see the RDT performed as expected in a variety of degraded network conditions.

Scenario	In File	Out File	Match	Time
Control	<i>dog.jpg</i>	<i>dog-control.jpg</i>	✓	0.002336s
Delay	<i>dog.jpg</i>	<i>dog-d.jpg</i>	✓	12.037570s
Loss	<i>dog.jpg</i>	<i>dog-l.jpg</i>	✓	57.010318s
Constrained Rate	<i>dog.jpg</i>	<i>dog-cr.jpg</i>	✓	0.930083s
Loss and Delay	<i>dog.jpg</i>	<i>dog-ld.jpg</i>	✓	92.969945
Loss and Constrained Rate	<i>dog.jpg</i>	<i>dog-lcr.jpg</i>	✓	64.960451s
Delay and Constrained Rate	<i>dog.jpg</i>	<i>dog-dcr.jpg</i>	✓	13.948315s
Delay, Loss and Constrained Rate	<i>dog.jpg</i>	<i>dog-ldcr.jpg</i>	✓	64.960451s

Table 1: Results for varying network conditions

4 Analysis

4.1 Bandwidth Delay Product and Link Utilisation

The Bandwidth-Delay Product (BDP) represents the ‘capacity’ of a link and is calculated as:

$$BDP = \text{Bandwidth (bits per second)} \times \text{Round Trip Time (seconds)}$$

For the link between pc7-082-1 and pc7-085-1, which was assumed to be 1Gbps Ethernet, the observed average RTT with RDT was 0.106ms. This yields a BDP of 13250 bytes for the link.

As RDT only sends 1312 bytes in each packet, this means that RDT achieves a link utilisation of 9.9%.

4.2 Maximum Theoretical Data Rate

We can also calculate that the maximum data rate of RDT as:

$$\text{Max Data Rate} = \frac{\text{Packet Size (bits)}}{\text{RTT (seconds)}} = \frac{1312 \times 8}{0.000106} = 98.9\text{Kbps} = 0.0989\text{Gbps}$$

4.3 Performance in Different Network Scenarios




Figure 5: Total Transmission Time in Different Network Scenarios for 230 Kb file.

From Figure ??, we can see that the network characteristic that has the biggest impact on the performance of RDT is packet loss. Packet delay has the second greatest effect on total transmission time, with constrained data seemingly having a minimal effect. This is most likely due to the RDT data rate being less than the constrained data rate for slurpe-3.

4.4 Efficiency of Packet Size

We can characterise the efficiency of the RDT packet size ρ as the ‘wastage due to control information’ [?], which for RDT is calculated as:

$$\rho = \frac{i - c}{i + a} = \frac{1312 - 12}{1312 + 12} = 0.982 \text{ (3 s.f.)}$$

5 Evaluation

5.1 Idle-RQ vs Continuous-RQ

The main weakness of RDT is its Idle-RQ mechanism. This leads to severe underutilisation of the data link, only approximately 9%. This is the inherent trade-off of the simple operation of Idle-RQ. Implementing Continuous-RQ with Go-Back-N would help to achieve higher link utilisation.

However, in doing so, the transmission rate becomes an issue. It is likely that full link utilisation may not be achievable if testing with `slurpe-3` due to limitations on the size of its internal queue. Resultingly, the choice of N for Go-Back-N would require careful consideration.

While link utilisation would improve, RDT with Go-Back-N would still be disproportionately impacted by dropped packets and the performance impact of re-transmission. Therefore, Selective re-transmission may yield better results.

5.2 Vulnerabilities

As with TCP, RDT provides no security mechanisms. As such, there are several vulnerabilities that may exist within RDT. Firstly, RDT is vulnerable Denial-Of-Service attacks via mechanisms such as SYN flooding. Secondly, the generation of Initial Sequence Numbers is pseudo-random and thus making ISM generation potentially predictable.

5.3 Extension Features

Given that RDT is implemented using UDP and uses the same checksum algorithm used by UDP, it is unlikely that RDT will ever encounter an incorrect checksum. This makes the use of a checksum in RDT mostly redundant. However, the inclusion of a checksum provided some utility. Firstly, it helped to catch several errors in the initial implementation of RDT during development, and secondly, it provided a useful opportunity to understand the function of the Internet Checksum [?].

However, the inclusion of adaptive RTO was more useful. From our analysis in Section ??, we concluded that the amount of re-transmission required has the biggest impact on RDT performance. Whilst RDT cannot do anything to minimise packet loss, it is able to control the number of re-transmissions due to RTO via the adaptive RTO mechanism. Using the adaptive RTOs, RDT can account for varying amounts of network delay.

Therefore we can say that adaptive RTO has had a positive impact on the efficiency and performance of RDT in scenarios where network delay is present.

5.4 Further Extension

While the implementation of bi-directional communication and Continuous-RQ was not implemented, it is useful to consider these features in evaluating the design of RDT.

In its current design, RDT would be unable to support simultaneous bi-directional communication due to its use of a two-way handshake. For bi-directional communication to occur, both parties are required to choose and synchronise an 'Initial Sequence Number', which is not possible with only a two-way handshake. Therefore, to support bi-directional communication, RDT would require a significant re-design.

However, RDT would not require a fundamental re-design to support Continuous-RQ. Continuous-RQ would solve RDT's fundamental issue of low transmission rate due to link under-utilisation (see Section

6 Conclusion

As demonstrated in this report, RDT provides reliable, ordered, uni-directional data transfer on top of UDP. RDT provides reliable performance in heavily degraded network environments. However, in its current implementation, RDT is extremely inefficient due to its Idle-RQ transmission mechanism. RDT performance suffers particularly when there is heavy packet loss that causes re-transmission. This efficiency could have been addressed, if given more time, by replacing Idle-RQ with Continuous-RQ with Go-Back-N.

Overall this practical provided an interesting opportunity to design and implement a protocol and analyse protocol performance.

7 Endnotes

1. Clock granularity is not considered, however RTO values are calculated in microseconds and the School Lab PCs have a granularity of 1 nanosecond.
2. The image was obtained under a royalty-free licence. Credit: Chayathon Wong Getty Images

A Appendix

A.1 Finite State Machine

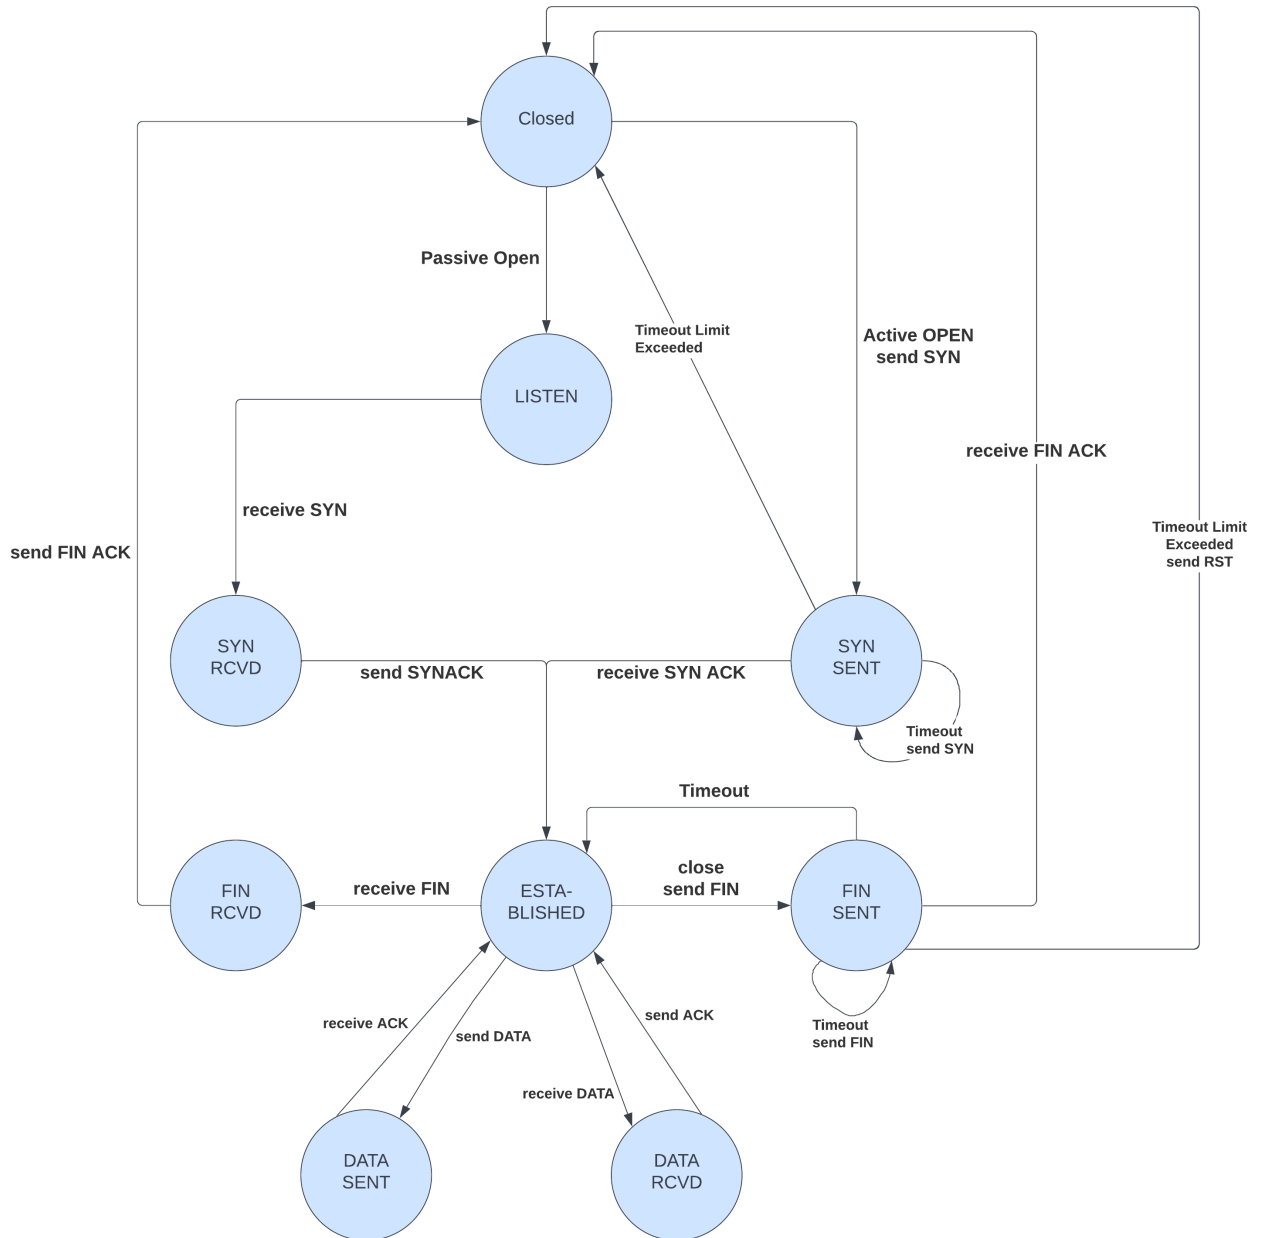


Figure 3: RDT Finite State Machine

A.2 Connection Management Timeline Diagram




Figure 4: Connection Timeline Diagram

A.3 Performance in Different Network Scenarios




Figure 5: Total Transmission Time in Different Network Scenarios for 230 Kb file.

A.4 Testing Screenshots

Testing screenshots can be found inside the `data/screenshots` directory.